

DANIEL KOMAROMY, EKOPARTY 2017

---

# EXPLORING AND BREAKING SAMSUNG'S TRUSTZONE SANDBOXES

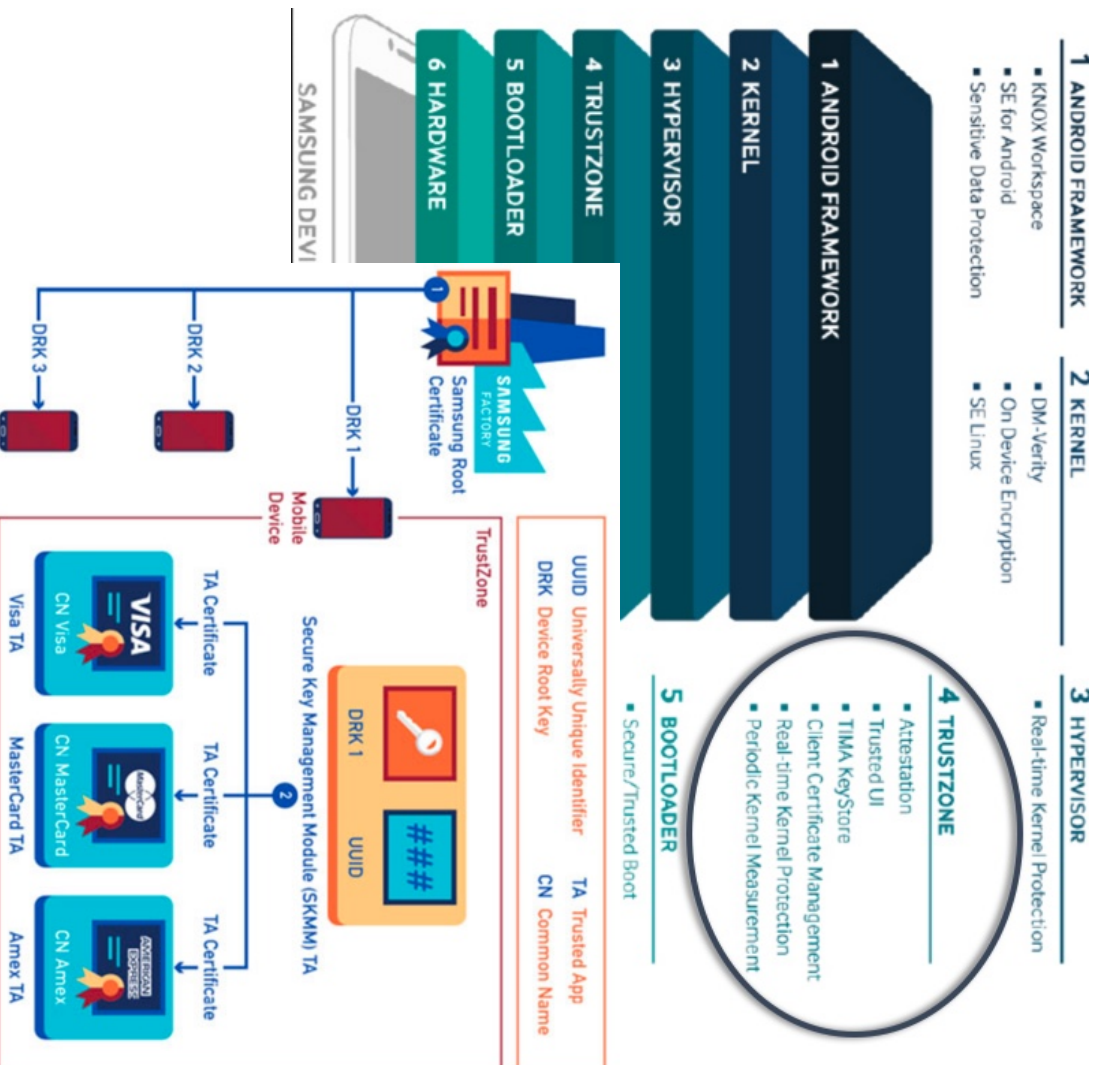
# INTRO

# WHY

- ▶ TrustZone on other platforms
  - ▶ extensively researched
  - ▶ statically found vulnerabilities, straightforward exploitation due to terrible isolation properties
- ▶ TrustZone on Samsung
  - ▶ minimal research, until recent work on Trustlets only (P0)
  - ▶ OS complete black box
- ▶ my main interest: understanding the system

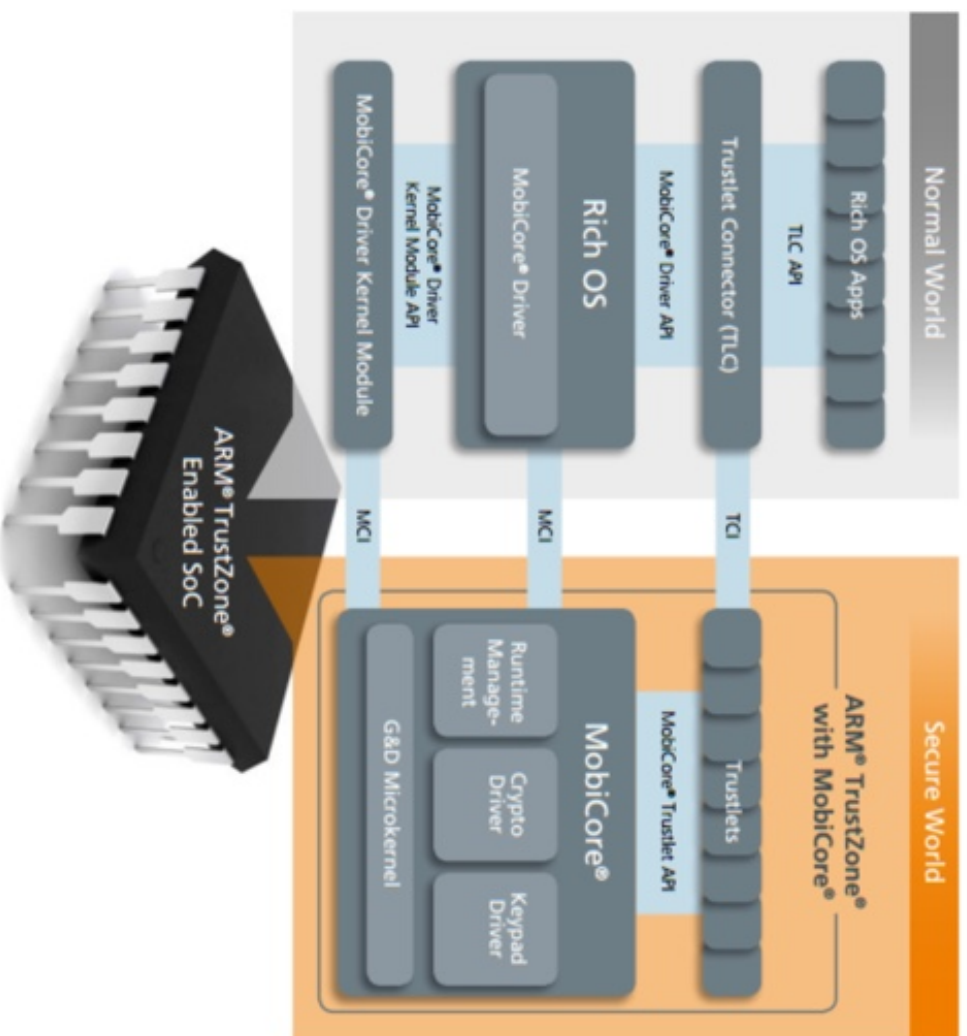
## INTRODUCTION

# WHAT WE KNOW ABOUT TRUSTZONE AND SAMSUNG KNOX



- ▶ Trusted Applications provide services: DRM, KeyStore, Certificate Management, SPay, Trusted UI, Trusted PIN, etc.
- ▶ Trusted Drivers control access to devices: SPI (Synaptic Fingerprint Sensor), eSE (NFC), MST (SPay), etc.
- ▶ Trusted Drivers add platform security to Android: Kernel Measurement, Device Encryption, etc.

# WHAT WE KNOW ABOUT THE TRUSTONIC TEE



- ▶ microkernel-based OS
- ▶ SW but NOT EL3 (monitor mode)
- ▶ MobiCore Control Protocol for nSW to talk to SW (documented)
- ▶ Trustonic provides Android {kernel driver, Daemon, client library} that proxy MCP commands into Android-level mcXYZ APIs
- ▶ Trustlet loading format (MCLF), names of APIs (tApiXYZ) provided by a common “libc” fairly well known / previously researched
- ▶ **Trustlet Connector Interface (TLC/TCL), aka the protocols for talking to Trustlets: common TCL header, otherwise entirely proprietary!**

# WHAT MAKES TBASE A MICRO KERNEL?

- ▶ handful of SMC fastcalls only, two that matter:
  - ▶ INIT: Android tells Trustonic location of command queues
  - ▶ NOTIFY: tell Trustonic there is a new command in the queues
- ▶ everything else is via shared memory, handled by:
  - ▶ MCP (loading/configuring trustlets): ??? but the expectation is that it is not the kernel itself
- ▶ Trusted Applications

# WHAT WE DON'T KNOW 1.

- ▶ How is the t-base microkernel implemented?
- ▶ Where can t-base firmware actually be found on Android?
- ▶ How does it start up and implement common features (MMU, process management, etc)
- ▶ What system calls does it implement?
- ▶ How/where are SMCs handled? Are there any extra SMC calls beyond what's documented?
- ▶ How/where is the MCP implemented?

# WHAT WE DON'T KNOW 2.

- ▶ Trustlets
  - ▶ How exactly are they loaded? E.g. how is the libc (mclib) loaded?
  - ▶ What Trustlet uid corresponds to what actual (KNOX) feature?
  - ▶ What are the differences between Trustlets and Drivers?
  - ▶ How are t!Api/drApi calls actually implemented, how do they map to interfaces towards the microkernel?
- ▶ Security hardening and isolation features in SW?



# WHAT WE DON'T KNOW 3.

- ▶ Android side
  - ▶ What happens beyond libMcClient.so?
  - ▶ What system components *actually* use TrustZone?
  - ▶ What debugging features are present?
  - ▶ What (if any) attack surface is there for an unprivileged user?

# REVERSING T-BASE

## FIRMWARE EXTRACTION

- ▶ t-base header embedded inside sbboot, identified by "t-base"
- ▶ gives addresses and sizes of several firmware components:
  - ▶ microkernel
  - ▶ mclib library
  - ▶ core trustlets (e.g. crypto driver)
  - ▶ the special "dom0" trustlet (MCCB/RTM/S0CB ... whatever :)

## FIRMWARE EXTRACTION

```

tbase_extract_table DCB "t-base", 0
;
; descriptor struct:
; char name[8]
; int offset
; int size
; char padding[0x10]
;
; real start offset: 0x132000
;
; Mtk: 0->0x147000 -> so that's the image itself -> so go back 0x15000 from "t-base"
; Image_h: 0x147000 -> 0x148000 -> so that's this
; Rtm: 0x148000 -> Rtm is SOCB (IPC/MCP implementing "demo" root trustlet)
;
; decrypt: 0x167000 -> MCLF header there
; tproxy: 0x17A000 -> MCLF header there
; sth2: 0x17B000 -> MCLF header there
; mclib: 0x183000 -> tlib indeed

DCB 0
DCD 0x5D000
ALIGN 0x20
DCB "mtk", 0
DCD 0
DCD 0x15000
ALIGN 0x20
DCB "image_h", 0
DCD 0x15000
DCD unk_1000
ALIGN 0x20
DCB "rtm", 0
DCD 0x16000
DCD 0x1P000
ALIGN 0x20
DCB "decrypt", 0
DCD 0x35000
DCD 0x13000
ALIGN 0x20
DCB "tproxy", 0
DCD 0x48000
DCD 0x1000
DCD 0
DCD 0
DCD 0
DCD 0
DCB "sth2", 0
DCD 0x49000
DCD unk_8000
ALIGN 0x20
DCB "mclib", 0
DCD 0x51000
DCD unk_C000
ALIGN 0x1000
DCB "SOCB_start", 0

```

## FIRMWARE ARCHITECTURE

- ▶ t-base is connected to nSW by ATF (ARM Trusted Firmware)
  - ▶ <https://blog.quarkslab.com/reverse-engineering-samsung-s6-sboot-part-ii.html> awesome explanation of this monitor mode, read that :)
- ▶ let's start from the microkernel initialization and identify: SMC handling (MC fastcall commands), MCP handling, SVC handling
- ▶ hopefully we'll understand what role the S0CB has

## T-BASE MICROKERNEL START UP

- ▶ loading to ida: start address 0x7F00000
  - ▶ can be figured out based on embedded pointer values, but also boot logs that can be googled. Prior work also mentioned this by now (Gal Beniamini)
- ▶ typical helpers: MCR instructions (setting VBAR, enabling MMU, etc)
- ▶ but in any case, it was still pretty painful. lot of things unclear at first that I was able to come back to when I understood more about higher layers

## T-BASE MICROKERNEL START UP

- ▶ early steps: pretty typical stuff
  - ▶ configure VBAR, enable MMU, enable cache, reset various context's stacks (interrupt, smc, etc)
- ▶ then starts the t-base specific part
  - ▶ tell ATF where the VBAR is for handling SMC calls
  - ▶ map the mclib and the S0CB
  - ▶ initialize process structures and start S0CB

# T-BASE MICROKERNEL START UP: SMC

- ▶ ATF needs to know the VBAR, because its SMC handler uses the offset from there to trap into t-base's smc/fastcall handler. From here we can find the SMC handlers, eg fastcalls:

```
tbase_smc_send_VBAR
LDR    R0, =0xB2000002
MOV    R1, #1
LDR    R2, =vector_table_normal
SMC    #0 ; SMC call tells ATF the VBAR address
BX     LR
; End of function tbase_smc_send_VBAR
```

```
void __fastcall __noreturn fastcall_handler(int a1, int a2, int a3, int a4, int a5)
{
    unsigned int main_id_reg; // r2@1

    main_id_reg = __mrc(15, 0, 0, 0, 5);
    __mcr(15, 0, vector_table_int_cxt, 12, 0, 0); // set the VBAR for fastcall handling
    if (cpu_revision_number == 4 * ((main_id_reg >> 8) & 0xF) + (main_id_reg & 0xF) )
        process_fastcall(a5, a2); // used here
    else
        handle_fastcall_2(a5, a2);
    __mcr(15, 0, vector_table_normal, 12, 0, 0); // restore normal VBAR
    __asm { SMC #0 } // back to ATF
    __isb();
    while ( 1 )
    ;
}
```

```
int __fastcall process_fastcall(int a1, int a2)
{
    int v2; // r4@1
    signed int *fc_args; // r0@1
    unsigned int fastcall_id_; // r1@1

    v2 = a2;
    fc_args = get_fastcall_args_addr(a1, a2);
    fastcall_id_ = *fc_args & 0xFF000000;
    if ( fastcall_id_ == 0x84000000 )
        return process_psci_fastcalls(fc_args, v2);
    if ( fastcall_id_ == 0xFF000000 )
        return process_mc_fastcalls(fc_args);
    return process_custom_fastcalls(fc_args, v2);
}
```



## T-BASE MICROKERNEL START UP: MAP LIB, SOCB

- ▶ to find where processes are mapped into memory, I had to find the functions used for managing page table entries
- ▶ identifying page table operations helped by:
  - ▶ MMU enablement code fixes TTBR0 to a static address, references easy to see
  - ▶ uses normal AARCH64 two level page table design, so code that e.g. creates a pte has a pretty easily recognizable look
  - ▶ intuition: it's readily apparent from the SMC handlers (fastcall and normal) that there isn't any implementation of MCP in here, so there HAD to be a mapping and creation of a new process in the init sequence (this turns out to be the SOCB).

## T-BASE MICROKERNEL START UP: REVERSING HINTS

- ▶ S0CB header format known, data flow from accesses to it in code can be followed
- ▶ matching syscalls made in similar operations in different places, e.g. syscall made by drApiMapClientBuffer to syscall made by mcMap MCP call, or functions used by syscall made to start process by mcOpenSession to functions used by microkernel startup sequence to start S0CB
- ▶ helpful: 0 heap in microkernel, all descriptors are at fixed locations, references trivial to see
- ▶ unlike in the microkernel, Samsung trustlets and drivers have a TON of debug strings; this helped a lot making sense of meaning of libc calls and from there figuring out what syscall implements what

## REVERSE ENGINEERING T-BASE

---

# T-BASE MICROKERNEL START UP: MAP LIB, SOCB

```
int __fastcall store_pte_0(int virt_addr, int virt_addr_high, int phys_addr_low, int phys_addr_high, int permissions)
{
    int virt_addr_; // r4e1
    int pt_entry; // r081

    virt_addr_ = virt_addr;
    pt_entry = create_pte(phys_addr_low, phys_addr_high, permissions, 0);
    return store_VA_to_pt(0, virt_addr_, pt_entry);
}

unsigned int __fastcall create_pte(int pa_low, int pa_hi, unsigned int perms, int dirtyflag)
{
    unsigned int pte; // r081
    unsigned int v5; // r483
    unsigned int v6; // r084

    pte_ = pa_low & 0xFFFFF000 | 0x403;
    if ( perms & 0x80 )
        pte_ |= 0x20u;
    v5 = (perms >> 1) & 1;
    if ( dirtyflag )
    {
        v6 = pte_ | 0x800;
        if ( v5 )
            pte_ = v6 | 0x40;
        else
            pte_ = v6 | 0xC0;
    }
    else if ( !v5 )
    {
        pte_ |= 0x80u;
    }
    if ( perms & 0x100 )
        return pte;
    if ( perms & 0x10 )
    {
        pte_ |= 4u;
    }
    else if ( perms & 8 )
    {
        pte_ |= 8u;
    }
    else
    {
        pte_ |= 0x30Cu;
    }
    return pte_;
}

void __fastcall __noreturn config_tbase_start_socb()
{
    process_struct_1 *new_proc; // r581
    unsigned int max_processes; // r181
    unsigned int v2; // r382
    int v3; // r084
    void *process_structs_area_addr; // [sp+0h] [bp-10h]@1
    int process_structs_area_length; // [sp+4h] [bp-Ch]@1

    memset_to_0(&tbase_conf_t, 60);
    tbase_conf_t.tbase_state = tbase_conf_t.tbase_state & 0xFFFFF8FF | 0x100;
    config_tbase_map_socb_to_0xFFFF0000_tell_atf_vbar(&process_structs_area_addr, &process_structs_area_length);
    sanity_check_SOCB_header();
    prepare_per_process_descriptors(
        process_structs_area_addr,
        process_structs_area_length,
        mapped_socb_max_processes,
        mapped_socb_max_sgs_allowed);
    dword_7F11780 = -1;
    write_EI1_thread_id();
    reset_counter_something();
    tbase_conf_t.tbase_state = tbase_conf_t.tbase_state & 0xFFFFF8FF | 0x200;
    new_proc = init_new_process(1, 0, 0xFFFF0F, (LOBYTE(mapped_socb_max_sgs_allowed) << 24) | 0xFFFFF);
    max_processes = mapped_socb_max_processes;
    while ( max_processes > 1 )
    {
        v2 = max_processes--;
        new_proc->perm_flags[v2 >> 5] |= 1 << (v2 & 0x1F);
    }
    v3 = mmap_socb_for_itself(new_proc);
    if ( create_new_process(new_proc, mapped_socb_entry, v3) )
    {
        _exit_to_ATF(15, 0x7F00976);
        while ( 1 )
        {
            // actually ENER into SOCB
            jump_to_R0(go_to_ATF_or_sw_EI0, 1, 0, 0);
        }
        _mcr(15, 0, EI1_thread_struct, 13, 0, 4);
    }
}
```

## T-BASE SYSCALLS

```

void __fastcall svc_D_mprotect(int r0_0)
{
    int v1; // r4@1
    int v2; // r1@1
    process_struct_1 *proc_t_map_into; // r4@3
    unsigned int arg1; // r9@4
    unsigned int arg2; // r7@4
    int arg3; // r10@4
    int arg4; // r8@4
    int v8; // r6@4
    process_struct_1 *proc_t_map_from; // r5@12
    int v10; // r0@18
    unsigned int id; // [sp+0h] [bp-10h]@1
    void *a1; // [sp+8h] [bp-8h]@1

    v1 = r0_0;
    a1 = r0_0;
    v2 = get_arg0(r0_0);
    id = get_arg0(a1) >> 16;
    if ( proc_id_check_1(id) && !proc_id_is_0(id) )
        proc_t_map_into = get_process_struct_by_id(id);
    else
        proc_t_map_into = *(v1 + 364);
    arg1 = get_arg1(a1);
    arg2 = get_arg2(a1);
    arg3 = get_arg3(a1);
    arg4 = get_arg4(a1);
    v8 = 1;
    if ( arg4 << 29
        && proc_id_check_1(v2)
        && (arg1 < arg2 || !arg2)
        && is_low_12_bits_0_address_alignment_check(arg1 | arg2 | arg3) )
    {
        if ( v2 )
        {
            proc_t_map_from = get_process_struct_by_id(v2);
            if ( (!proc_id_check_1(id) || proc_id_is_0(id)) && proc_t_map_from != proc_t_map_into )
            {
                v8 = verify_process_permission(proc_t_map_from, v2, proc_t_map_into);
                if ( v8 )
                    goto LABEL_19;
            }
        }
        else
        {
            proc_t_map_from = proc_t_map_into;
        }
    }
    execute_mprotect(proc_t_map_from, proc_t_map_into, arg1, arg2, arg3, arg4);
}

//
// from the process struct, get the first level page table,
// from there get the phys address of corresponding second
// level page table and map that into a fix address for the kernel.
// then, now sitting in kernel virtual address, directly modify
// that page table entry to actually create the mapping.
//
syscall_table
DCD svc_0_nop+1 ; DATA XREF: invoke_syscall_from_table+4070 ; invoke_syscall_from_table:syscall_table_ptr0
DCD svc_1_init_process+1
DCD svc_2_nop+1
DCD svc_3_nop+1
DCD svc_4+1 ; Did not find this invoked anywhere in {SOCB,tlib}
DCD svc_5_start_process+1
DCD svc_exit+1
DCD svc_mmap+1
DCD svc_8_munmap+1
DCD svc_9_start_thread+1
DCD svc_A_stop_thread+1
DCD svc_B_return_0xd+1
DCD svc_C_modify_thread_registers+1
DCD svc_D_mprotect+1
DCD svc_E_resume_thread+1
DCD svc_F+1
DCD svc_10_set_thread_prio+1
DCD svc_11_ipc+1
DCD svc_12_int_attach+1
DCD svc_13_int_detach+1
DCD svc_14_signal+1
DCD svc_15_signal+1
DCD svc_16+1 ; Did not find this invoked anywhere in {SOCB,tlib}
DCD svc_tbase_svc_fastcall_input+1
DCD svc_18_log_char+1
DCD svc_19_get_secure_timestamp+1
DCD svc_1a_control+1 ; includes a lot, such as:
; - driver shmem map/unmap
; - get/set exception info
; - get MCP queue info
; - get IPCM phys address values
; - cache control
; - virt2phys, phys2virt translation
; - set custom fastcall, call custom fastcall

```

# REVERSE ENGINEERING SOCB

- ▶ SOCB is a trustlet also, with (almost) same MCLF loading format
- ▶ Start-up sequence is pretty complex, as it has to initialize various descriptors for maintaining state information of trustlet instances and it also has to load initial drivers (specifically, the crypto driver)
- ▶ The most important is starting three threads:
  - ▶ MCP handling
  - ▶ IPC handling
  - ▶ Trustlet notification handling
- ▶ To be able to handle things, SOCB needs to access the MCP and notification queues. It uses the control syscall (0x1A) to get the address information from the microkernel and then maps it in for itself.



## REVERSE ENGINEERING SOCB: IPC

- ▶ IPC: entirely proprietary, no initial info.
- ▶ The big help was identifying the tlApi/drApi callers of SVC 0x11, which gave away which command type is what.
- ▶ From there, connect more syscall knowledge (e.g. mmap, mprotect etc) to commands
- ▶ It was an iterative process, constantly back and forth among microkernel, SOCB, and trustlets
- ▶ In total: of 40 IPC commands, I was able to identify the purpose of 30



## REVERSE ENGINEERING SOCB: IPC

```

case 1u:
    ret.field_0 = ipc_cmd;
    ret.msg_len = 0;
    goto PROCESS_REQ;

case 2u:
    MCP_process_ipc_0x2_MSG_RESP(peerIDs.caller_ID, ipc_cmd);
    break;

case 3u:
    ret.field_0 = ipc_arg;

    // MSG_REQ -> Request from client to server (ta->driver)
    // MSG_RS -> Response
    MCP_process_ipc_peerIDs.called_ID, ipc_cmd);

    break;

    // MSG_RD -> Driver Ready: first IPCR call by a driver always
    // basically the very first time we call from a driver,
    // we will get assigned a (MSG) called ID, this is essentially
    // the driver ID that is the same for each instance.
    ret.msg_len = peerIDs.caller_ID;
    *a2 = ipc_cmd;
    a4 = peerIDs.called_ID;
    mcp_drv_set_id_2(driverId(ret.msg_len, ret.msg_len));
    TA_instance_set_id_8(driverId(ret.msg_len, ret.msg_len));
    *v1 = find_MCP_DRV_struct_ptr(ret.msg_len);
    if (drv)
    {
        marshalled_cmd.ta_id = a4;

        // here the ipc_cmd parsed struct & marshalled cmd
        // struct is actually var_84, so that's really confusing
        v30 = *a2;
        *marshalled_cmd.handler_id = ret;
        call_handler_by_id(drv, v30);
        // we are going to call the handler 3 here
        call_MSG_HANDLER_of_all_TMS();
    }
    break;

case 4u:
    v2 = find_trustlet_by_driverId(peerIDs.caller_ID);
    if (v2 != -1)
    {
        // MSG_NOTIFY
        do_notify_Ox1a_control_cmd(v2);
    }
    break;

case 5u:
    peerIDs.called_ID;
    ret = *(v30 - 1);
    // MSG_CLOSE_TRUSTLET
    *marshalled_cmd.ta_id = ret;
    marshalled_cmd_ipc_arg = peerIDs.caller_ID;
    if ( sub_2B56(4, peerIDs.caller_ID) )
    {
        if ( marshalled_cmd.ta_id != -1 )
        {
            v30 = marshalled_cmd_ipc_arg;
            *marshalled_cmd.handler_id = ret.field_0;
            *a2 = ret.msg_len;
            a4 = ret.field_0;
            v6 = find_mcp_TA_struct(ret.field_0);
            if ( v6 )
            {
                call_handler_by_id(v6, a2);
                dword_P22C = sub_319C(a2);
                if ( dword_P22C )
                {
                    sub_2460();
                }
            }
        }
    }
    break;
}

// MSG_NOTIFY_HANDLER

```

```

case 11u:
    // MSG_GET_DRIVER_VERSION
    ret = peerIDs;
    if ( TA_with_id_loaded(peerIDs.caller_ID) )
    {
        *a2 = 0;
        if ( ret.msg_len )
        {
            v12 = mcp_drv_find_by_id_unknown_3(ret.msg_len);
            if ( v12 )
            {
                *a2 = v12->driver_version;
            }
        }
        MCP_process_ipc_cmd_complete_response(ret.field_0, 0, 11, *a2);
    }
    break;

case 12u:
    // MSG_GET_DRAPI_VERSION
    ret.field_0 = peerIDs.caller_ID;
    memset_to_0(a2, 96);
    if ( sub_3632(ret.field_0) & 1)
    {
        MCP_process_ipc_cmd_complete_response(ret.field_0, 0, 12, resp_data);
    }
    break;

case 13u:
    // MSG_SEM_NOTIFICATION_HANDLER
    ret.field_0 = peerIDs.called_ID;
    ret.msg_len = peerIDs.caller_ID;
    if ( is_trustlet_type_a_driver_SECURITY_CHECK(peerIDs.caller_ID) || is_trustlet_systemTrustlet(ret.msg_len) )
    {
        a4 LOWORD(ret.msg_len) == LOWORD(ret.field_0) )
        {
            sub_8394(ret.msg_len, ret.field_0);
            MCP_process_ipc_cmd_complete_response(ret.msg_len, 0, 13, ret.field_0);
        }
    }
    break;

case 14u:
    // MSG_GET_REGISTER_ENTRY
    MCP_process_ipc_0x4(peerIDs.caller_ID, ipc_arg);
    break;

case 15u:
    // MSG_DRV_NOT_A_DRAPI_NOT_CLIENT
    ret = 0;
    *a2 = 0;
    if ( is_trustlet_type_a_driver_SECURITY_CHECK(peerIDs.caller_ID) || is_trustlet_systemTrustlet(ret.field_0) )
    {
        if ( TA_with_id_loaded(ret.msg_len) & (wrap_find_trustlet_instance_by_driverId(ret.msg_len), v16 == 2) )
        {
            call_ipc_signal(ret.msg_len);
            *a2 = 3847;
        }
        else
        {
            *a2 = 3843;
        }
    }
    MCP_process_ipc_cmd_complete_response(ret.field_0, 0, 2, *a2);
    break;

case 16u:
    // MSG_SEM_PASTCALL_HANDLER
    // (dramInstallIPC)
    ret.msg_len = peerIDs.called_ID;
    a4 = ipc_arg;
    *marshalled_cmd.handler_id = 0;
    if ( is_trustlet_type_a_driver_SECURITY_CHECK(peerIDs.caller_ID) )
    {
        if ( *a2 >= 0x100000u )
        {
            *marshalled_cmd.handler_id = 3842;
            v17 = find_trustlet_by_driverId(a4);
            if ( sub_7f9a(v17) )
            {
                v18 = syscall_Ox1a_0x94, 0, a4, *a2, ret.msg_len, 0); // SEM_CUSTOM_PASTCALL_HANDLER
                if ( v18 )
                {
                    *marshalled_cmd.handler_id = ((v18 & 0x7ff) << 12) + 3848;
                }
            }
            else
            {
                *marshalled_cmd.handler_id = 3848;
            }
        }
        else
        {
            *marshalled_cmd.handler_id = 3848;
        }
    }
    else
    {
        *marshalled_cmd.handler_id = 3848;
    }
}

```



# REVERSE ENGINEERING SOCB: TRUSTLET NOTIFY

```
void __noreturn attach_intr_init_mcp_queues_and_process_notifications_forever()
{
    DWORD *v4; // r0e1
    int *v5; // r0e3

    __asm { SVC          0x12 }

    v4 = custom_log_reserve_sg(4);
    *v4 = 0x60100001;
    assert(v4, &custom_heap_something, 0);
    while ( 1 )
    {
        call_init_MCP_queues();

        // uses a control syscall (0x1A) to get the queue addresses
        // that were set up when nSW called INIT fastcall,
        // then it mmap's the queues into its own address space.
        //

        do
        {
            while ( 1 )
            {
                v5 = get_noti_from_notification_queue();
                if ( !v5 )
                    break;
                send_signal_to_ta(*v5);

                // svc 0x15
                svc_0x11_to_0x8001PFPF();
            }
            while ( dword_18B50 );
            maps_or_unmaps_mcp();
        }
    }
}
```

# TRUSTLETS

# TRUSTLETS AND DRIVERS

- ▶ similarities and differences
- ▶ all depends on tlApi and drApi calls
- ▶ Trustlet behavior: single threaded; input from WSM tcIBuf, tlApi\_callDriver if driver needed
- ▶ Driver behavior: often multi threaded (handle IRQs from devices)
  - ▶ input from drApiIpccCallToIPCH: command type + length + initial command pointer
  - ▶ passed in arguments (marshalled from tlApi\_callDriver by S0CB)
  - ▶ then map in Trustlet memory to read actual command buffer(s) using more drApi calls that map to S0CB IPC commands
- ▶ so, the comm. b/w trustlets&drivers is marshalled by the previously seen S0CB
  - ▶ *no* "MAC" on IPCs at all; up to drivers to filter callers for commands based on caller uid

## LABELING TLLIB APIS

- ▶ Prior art + strings in Samsung TAs covered most tLApis, but few drApis; additional reversing of secure drivers helped fill in the gaps
- ▶ “GOT” approach is really trivial: one field in trustlet/driver headers gets filled in with tLib address at load time; all library calls simply call to this address with R0 holding the API id
- ▶ Luckily, all such calls go through small stub functions, so we can label these nicely in the code.
- ▶ Wrote a script to automatically label the APIs in a trustlet/driver binary, to aid reverse engineering



# TRUSTLETS AND DRIVERS

- ▶ 0006...: Widewine trustlet (drm)
- ▶ 0701...: TlCm (certificate management) trustlet
- ▶ 0706...: TEE keymaster trustlet
- ▶ 0813...: TEE gatekeeper trustlet
- ▶ f...5: HDCP trustlet (drm)
- ▶ f...a: TIMA PKM trustlet
- ▶ f...b: TIMA LKM Auth trustlet
- ▶ f...c: Key Management trustlet
- ▶ f...d: ?
- ▶ f...e: Synaptics Fingerprint trustlet
- ▶ f...f: TIMA attestation trustlet

# TRUSTLETS AND DRIVERS

- ▶ f...12: CCM trustlet
- ▶ f...13: Keystore trustlet
- ▶ f...14: TUI trustlet
- ▶ f...16: SKMM trustlet (Secure Key Management Module, used by SPay trustlets)
- ▶ f...17: MLDAP trustlet (knox mdm integration <https://www.samsungknox.com/en/solutions/knox-workspace/on-premise-mdm-integration>)
- ▶ f...19: Dmverity trustlet
- ▶ f...1e: OTP trustlet
- ▶ f...1f: FIDO trustlet
- ▶ f...2e: Fingerprint trustlet
- ▶ f...38: ESECOMM trustlet
- ▶ f...3e: TEE Keymaster
- ▶ f...41: ICC

# TRUSTLETS AND DRIVERS

- ▶ f...1c: VISA Pay
- ▶ f...21: Mastercard Pay
- ▶ f...27: PLCC Pay
- ▶ f...28: KRCC Pay
- ▶ f...31: Discovery Pay
- ▶ f...3a: CHNCMM Pay
- ▶ f...33: JIC Pay
- ▶ f...32: CNCC Pay
- ▶ f...39: EURCOMM Pay



## TRUSTLETS AND DRIVERS

- ▶ ffffff0....1: SRPMB driver
- ▶ ffffffd....4: Crypto driver
- ▶ ffffffd....a: TIMA driver
- ▶ ffffffd....e: SEC SPI driver #1 (fingerprint sensor)
- ▶ ffffffd....14: TUI driver
- ▶ ffffffd....17: SEC SPI driver #2 (eSE)

# TLC IN ANDROID

## T-BASE AND ANDROID

- ▶ Debugging info in filesystem?

- ▶ /system/bin/tima\_dump\_log
  - /system/tima\_measurement\_info (attestation)
  - /system/kern\_sec\_info{1|2|3|4} (pkm)
  - /proc/tima\_secure\_log
  - /proc/tima\_secure\_rkp\_log
  - /proc/tima\_debug\_log
  - /proc/tima\_debug\_rkp\_log (rkp)
  - /sys/kernel/debug/trustonic\_tee/\*
- /proc/sec\_log**

## T-BASE AND ANDROID

- ▶ Looking around
  - ▶ `strings -f * | grep onTransact`  
`strings -f * | grep mcNotify`
  - ▶ lib names in `/system/lib64/` & `/system/vendor/lib64` with `"tlibc"`
  - ▶ `service list | grep {com.sec,finger,keystore} etc.`
  - ▶ finally, settled on `tlc_server`. Exposes access to certain trustlets via Binder. (ESECOMM, CCM, DCM, TUI)

## TLC SERVER

- ▶ one instance started for each TA (CCM, DCM, TUI, ESECOMM, PUF)
- ▶ uses dlopen/dlsym to get comm\_data initialization function, e.g. esecomm\_get\_comm\_data in libtcl\_tz\_esecomm.so; this identifies trustlet to load
- ▶ registers binder service, handler a simple switch of commands: OPEN, CLOSE, COMM, COMM\_VIA\_ASHMEM

# TLC SERVER

- ▶ OPENSWCONN: create\_comm\_cxt()
- ▶ COMM, COMM\_VIA\_ASHMEM: parse arguments out of Binder parcel and send via libMcClient.so
- ▶ COMM: uses SEAMS to validate caller based on SEAndroid policy
- ▶ COMM\_VIA\_ASHMEM: no authentication! **COMPLETELY OPEN**

```
comm_cxt = create_comm_cxt(  
  
    // TLC_COMM_TYPE:  
    // 0 - proxy  
    // 1 - direct  
    //  
    // ==> we create direct  
    //  
    // direct comm_cxt()  
    // - instantiate directCommImpl with these parameters  
    //   - includes root (== device id, switched out to 0) and process (==uid)  
    //   - call tlc open  
    //   - mcOpenDevice(0) and mcOpenSession to uuid,  
    // after it maps the TLC buffer, to sendmsglen+recvmsglen length  
    TLC_COMMUNICATION_TYPE_DIRECT,  
    comm_data_root,  
    comm_data_root_strlen,  
    comm_data_process,  
    comm_data_process_strlen,  
    comm_data_max_sendmsg_size,  
    comm_data_max_recvmsg_size);
```



# VULNERABILITIES

# ATTACK SURFACES

- ▶ unprivileged Android > TLC
- ▶ privileged Android > Trustlet (tci)
- ▶ privileged Android > IPCH (MCP implementation)
- ▶ privileged Android > tbase kernel (SMC)
- ▶ privileged Android > custom tbase fastcalls (SMC) and ATF SMCs
- ▶ Trustlet > Driver (IPC)
- ▶ Trustlet > tbase kernel (SVC)
- ▶ Trustlet > IPCH (IPC)



## FOUND VULNERABILITIES

- ▶ Authentication bypass in tlc\_server
- ▶ Memory corruption in tlc\_server
- ▶ Plenty of memory corruption vulnerabilities in trustlets (buffer overflows, arbitrary writes, integer overflow-to-buffer overflows)
- ▶ Session hijack logic vulnerability in CCM Trustlet
- ▶ Memory corruption and race condition vulnerabilities in TIMA driver
- ▶ Feature abuse in TIMA driver leading to arbitrary Android kernel read (KASLR bypass)

# STACK BOF IN ESECOMM

- ▶ ESECOMM: implements interface to eSE (NFC enabled for Samsung Pay)
- ▶ smart card communication: ISO7816
  - ▶ [http://www.cardwerk.com/smartcards/smartcard\\_standard\\_ISO7816-4\\_4\\_abrev\\_and\\_notation.aspx](http://www.cardwerk.com/smartcards/smartcard_standard_ISO7816-4_4_abrev_and_notation.aspx)
  - ▶ Uses SEC SPI driver to talk to eSE
  - ▶ SPI driver: there is also a Linux kernel driver (see /drivers/spi/spi-s3c64xx.c), this explains meaning of memory mapped I/O
- ▶ Implements "SCP03 Global Platform Secure Channel Protocol"
  - ▶ <http://csrc.nist.gov/groups/ST/ssr2016/documents/presentation-mon-traore.pdf>
  - ▶ uses an APDU based protocol with TLV parameters to set up secure channel cryptographic information (Diffie-Hellman keys)

# STACK BOF IN ESECCOMM

```
signed int __fastcall process_ScpiInstallCaCert(tci_msg_add_ca_payload_t *reqmsg, tci_rsp_payload_t *rspmsg)
{
    tci_msg_add_ca_payload_t *reqmsg_; // r4@1
    tci_rsp_payload_t *rspmsg_; // r5@1
    char *reqmsg_payload_; // r7@1
    signed int result; // r0@2

    char pubkey[512]; // [sp+Ch] [bp-244h]@2
    char caid[32]; // [sp+20Ch] [bp-44h]@2
    char curveid; // [sp+22Ch] [bp-24h]@2
    int pubkey_len; // [sp+230h] [bp-20h]@2
    unsigned int caid_len; // [sp+234h] [bp-1Ch]@2

    reqmsg_ = reqmsg;
    rspmsg_ = rspmsg;
    reqmsg_payload_ = reqmsg->payload;
    if ( validate_input_len(reqmsg->payload, reqmsg->total_len, reqmsg->payload, (char *)&reqmsg->total_len) ) //
        ONLY verifies the total input len, not related to TLV lengths
    {
        result = parse_ca_cert(reqmsg_payload_, reqmsg->total_len, caid, (int *)&caid_len, &curveid, pubkey,
            &pubkey_len);
        (...)
    }
```

# STACK BOF IN ESECOMM

```
signed int __fastcall parse_ca_cert(char *msg_payload, int total_req_payload_len, void *caid, int *caid_len, char *curveid,
void *pubkey_buf, int *pubkey_buf_len)
{
    (...)
    v7 = msg_payload;
    total_req_payload_len_ = total_req_payload_len;
    caid_ = caid;
    v10 = caid_len;
    v11 = -1;
    memset_to_0(&out_buf, 0x44u);
    if ( parse_tlvs_from_APDU(&out_buf, v7, 0, total_req_payload_len_) < 0 )
    {
        snprintf(logbuffer, 119, "%s:%d :: Error, %s\n", "parse_ca_cert", 73, "failed to parse TLV");
        logbuffer[119] = 0;
        tIApiLogPrintf_0("%s\n", logbuffer);
        return 3;
    }
    tlv_caid = find_tlv_obj_in_parsed_tlvs(&out_buf, (char *)&caid_tag_value);
    tlv_caid_ = tlv_caid;
    if ( tlv_caid )
    {
        memcpy_w(caid_, &tlv_caid->tlv_value_OR_num_extra_tlvs, tlv_caid->tlv_length);
    }
}
```

## CCM SESSION HIJACK

- ▶ CCM: Client Certificate Management
  - ▶ supposed to secure private certs, basically
    - ▶ [https://seap.samsung.com/api-references/android-premium/reference/com/sec/enterprise/knox/ccm/TZ\\_CCM\\_PKCS11\\_Guide.pdf](https://seap.samsung.com/api-references/android-premium/reference/com/sec/enterprise/knox/ccm/TZ_CCM_PKCS11_Guide.pdf)
- ▶ Keys never leave TZ, CCM is an encryption/decryption oracle.
- ▶ To secure this, at CA install time certs are secured with some authentication method (password, TUI)

## CCM SESSION HIJACK

- ▶ So, how does login exactly work?
- ▶ First, session must be opened with C\_OpenSession
  - ▶ 64 bit random session id generated
- ▶ Then, session must login, using C\_Login
  - ▶ pass session id to identify session
  - ▶ if successful, session->logged\_in bit flipped
- ▶ After this, until C\_CloseSession, all operations will be allowed where session id matches

## CCM SESSION HIJACK

- ▶ That creates a race condition window:
  - ▶ If legitimate user logs a session in,
  - ▶ ANYBODY can impersonate the session that knows the session id
- ▶ 64 bit randomness from crypto engine though, so that seems solid
- ▶ Except...

## CCM SESSION HIJACK

- ▶ The session ids are printed into the /proc/sec\_log ...
- ▶ Even if an app context would be restricted: the actual secure log implementation is the same buffer for ALL trustlets: by hijacking ESECOMM, we can read the sec\_log even if we can't read /proc/sec\_log itself b/c of SEAndroid

```
v16 = TZ_get_rand_data(rand, &a2);  
if ( v16 || a2 != 8 )  
{  
    snprintf(logbuff, 119, "TL_TZ_CCM: C_OpenSession rand gen failed with ret = 0x%x", v16, v20, v21, v22);  
    logbuff[119] = 0;  
    tlapilogvPrintf2("%s\n", logbuff);  
    v4 = 5;  
}  
else  
{  
    snprintf(logbuff, 119, "TL_TZ_CCM: OpenSession ID: 0x%x, len: %d", 0, *(_DWORD *)rand, v25, a2);  
    logbuff[119] = 0;  
    tlapilogvPrintf2("%s\n", logbuff);  
}
```



## TIMA DRIVER KASLR BYPASS

- ▶ Once we hijack ESECOMM, we can talk to the TIMA driver
- ▶ Implements many different things:
  - ▶ TIMA measurement golden copy area read/write (trustlets use this to verify/set warranty info)
  - ▶ nSW kernel hashing (for TIMA PKM) (every UUID!)
  - ▶ MST driver (for Samsung Pay) (only whitelisted UUIDs)
  - ▶ SCrypto (a full openssl stack based FIPS compliant crypto library... in addition to the Crypto Driver, hmmm.) (every UUID!)
- ▶ modify UUID whitelists (restricted to CNCC trustlet)

## TIMA DRIVER KASLR BYPASS

- ▶ measurement info read/write:
- ▶ this would allow arbitrary physical address read/write
- ▶ Gal Beniamini described this
- ▶ however, in practice, range checks limit what can be written to the intended areas
- ▶ Nonetheless - you could STILL replace the golden measurement/warranty info (don't try this at home...)

# TIMA DRIVER KASLR BYPASS

```
{ signed int __fastcall is_addr_range_valid_for_phys_access(unsigned int addr_start, int len, int access_type)
{
    unsigned int addr_end; // r301
    const char *v4; // r202
    unsigned int tima_unk_phys_addr_end; // r003
    int v6; // r3015
    const char *v7; // r2015

    addr_end = addr_start + len;
    if ( addr_start + len >= addr_start )
    {
        tima_unk_phys_addr_end = addr_from_tima_sfr_start_of_SW + 0x100000;
        if ( access_type )
        {
            if ( access_type != 1 )
            {
                v4 = "OMG!!! non-supported action";
                goto LABEL_6;
            }
            if ( addr_start >= addr_from_tima_sfr_start_of_SW + 0x100000 && addr_end <= tima_unk_phys_addr_end
                || addr_start >= addr_from_tima_sfr_start_of_SW + 0x164 && addr_end <= addr_from_tima_sfr_start_of_SW + 0x360 )
            {
                // all right, no quite clearly we are NOT
                // allowed to just write to any physical address,
                // unfortunately, because we are limited to this
                // TIMA range.
            }
        }
        return 0;
    }
    v6 = addr_start;
    v7 = "OMG!!! physical address %x is not allowed for writing";
    else
    {
        if ( addr_start >= addr_from_tima_sfr_start_of_SW && addr_end <= tima_unk_phys_addr_end
            || addr_start >= 0x80000000 && addr_end <= 0x80200000
            || addr_start >= 0x2023800 && addr_end <= 0x2024000
            || addr_start >= 0x80000000 && addr_end <= addr_from_tima_sfr_start_of_SW )
        {
            return 0;
        }
        v6 = addr_start;
        v7 = "OMG!!! physical address %x is not allowed for reading";
    }
    snprintf((int)logbuff, 119, v7, v6);
    goto LABEL_23;
}
v4 = "addr_pa + len overflow";
LABEL_6:
    snprintf((int)logbuff, 119, v4, addr_end);
LABEL_23:
    logbuff[119] = 0;
    debuglogvprintf_Wrep("%s\n", logbuff);
    return 1;
}
```

# TIMA DRIVER KASLR BYPASS

- ▶ What's left in TIMA?
  - ▶ SCrypto: I actually found several memory corruption vulnerabilities in here. Disclosure not finished yet.
  - ▶ PKM MD5 hashing command
- ▶ Hash command is open to all UUIDs! (Should have been restricted to TIMA PKM trustlet)
- ▶ address range allows all of physical memory within which nSW lives
  - ▶ to break Android KASLR: invoke from hijacked ESECOMM to hash first page of kernel physical address space: this is an empty page with just the value of the KASLR slide in it
- ▶ from the hash, it is a <32 bit bruteforce to recover the slide.

# TAKEAWAYS

- ▶ even an almost entirely black box TEE OS implementation can be reversed fairly comprehensively
- ▶ t-base architecture solid for isolation, but lacks modern exploit mitigations
- ▶ driver api design less susceptible to exploitability of input parsing memory corruption vulns, much worse for trustlets
  - ▶ still, architecture should find a way to sandbox secure drivers, e.g. filter what areas can be mapped by drivers, what drivers can make fastcalls etc.
  - ▶ improve driver whitelists for allowed commands!
- ▶ altogether, t-base provides a huge attack surface to an elevated Android privilege context and it does not help vendors to eliminate or deal with unprivileged Android attack surface

QUESTIONS?